Part A codebook, version 2 (issued 28 October 2026)

Supports §5.2, §6 and Appendix D. Version 2 changes version 1 only in F3a (place) and F3b (time), which now state when `n/a` applies; F3c (criterion) is unchanged. The second coder re-coded F3a–F3c of the sampled items under this version, and all agreement the article reports for these three fields comes from that re-code (§5.2, Appendix E). The message to the second coder that the status note mentions is not included. As for version 1, an AI model proposed each code and the first author reviewed every proposal (§5.5); the "Procedure" section is the plan written before coding (Appendix D).

Edits for anonymity: the first author's name became "the first author" (3 places), the building's name "Building A" (2 places), and "His" (the second coder) became "The second coder's" (2 places). SHA-256 of the original, unedited file: `0dda2b06ad8f3de187df61c66293f43ca07af5116b65dfacc5dda1cdcc470beb`.

Part A codebook — v2 (tracker F1)

Status: v2, issued 28 October 2026. Supersedes v1 (frozen 24 October 2026, SHA-256 in `CODEBOOK_PartA.sha256`, kept unchanged as the record of what the second coder was actually given). **Changed from v1:** F3a and F3b now state an explicit `n/a` rule; F3c's own definition is untouched. Found on 28 October, when the second coder's F3 codes were checked against the first author's: across all 160 sampled items (100 Part A + 60 mixed), the second coder never once used `n/a` for F3a or F3b, although v1 only ever defined `n/a` for F3c (one worked example) and said nothing about when a place or a time does not apply at all. That gap is fixed below. The SHA-256 of this file is in `CODEBOOK_PartA_v2.sha256`. From now on it changes only by issuing v3 with a new hash, never by editing.

F3c is a separate, larger problem that this patch does not fix. On the same 160 items the first author called `specified` on 9% of the Part A sample and the second coder called it `specified` on 72% — the opposite of a missing category; it looks like a real difference in what counts as "a threshold, band or comparison standard is stated". That has been put to the second coder directly (see the message below) rather than resolved here. The second coder's F3a/F3b/F3c codes on the 160 sampled items are being re-collected under this version, blind to the first author's codes and to the disagreement counts, before agreement on those three fields is reported. The second coder's F1, F2, F4, F5 and F6 codes are unaffected and already final.

Unit. One archived submission row: the question text, and any context the archive recorded with it (stage, scenario prompt where shown). **Never shown to coders:** the persona, the automated labels (complexity tier, lexicon class, domain), or the other coder's codes.

Examples below are constructed, not taken from the corpus (the corpus may not be quoted without permission checks).

F1 Request type (multi-label)

What operation does the asker need performed? Code every operation the wording **requires**, not every one it permits.

Value	Definition	Constructed example
`fact`	Look up one current or stated value or attribute	"What temperature is the lab?"
`compare`	Two or more named things or periods against each other	"Is floor 2 warmer than floor 3?"
`summarise`	Aggregate over a period or set (average, total, trend)	"How much energy did we use last week?"
`rank`	Choose or order among candidates by a criterion	"Which meeting room is quietest?"

Value	Definition	Constructed example
`explain`	Why something happened; a cause	"Why is it so stuffy in here?"
`capability`	Whether the building has or offers something; availability	"Is there a quiet room I can use?"
`action`	Asks the system to do something in the world	"Turn the heating up."
`unclear`	No operation can be determined	"Building?"

Rule. "Is X too hot?" is `fact` (with a criterion, F3), not `explain`. "Can I book X?" is `capability` + `action` only if it asks the system to make the booking.

F2 Evidence source required (multi-label)

Where must a **complete** answer come from? Code the source the answer **needs**, not the source the asker mentions.

Value	Definition	Constructed example
`measurement`	Sensor time series (readings over time)	"Was the room above 25 °C this afternoon?"
`description`	Static building-model facts: room type, capacity, equipment, location, layout	"How many people does room 2.13 hold?"
`record`	Dated operational entries: bookings, services, incidents, tests, work orders	"When was the fire alarm last tested?"
`document`	Policies, procedures, manuals, rules	"What do I do if the alarm sounds?"
`external`	Anything outside the building's own data: weather, transport, the web, people's opinions	"Will it rain at lunchtime?"
`none`	No evidence needed (greeting, opinion request) or cannot tell	"Hello"

Rule. Code `record` + `description` for "Is room 2.13 free now?" (bookings and room identity), and add `measurement` only if the wording asks whether anyone is *there*.

F3 Specification (three sub-codes, each `specified` / `missing` / `n/a`)

Sub-code	`specified` when...	`n/a` when...	Notes
F3a place	a room, floor, zone, asset or "the building" is named	**(v2)** the question needs no place at all — F2 is `none` ("What is OntoSage?"), or it asks about the building as a single whole with no finer place that could ever apply ("What is the building's total floor area?")	"here", "this room" = `missing` (the deictic needs context the text lacks). **Do not infer a place from anything outside the row**
F3b time	a time, date, period or "now/today" is stated	**(v2)** the answer is a fact that does not vary with time — capacity, room type, policy, a fixed equipment list ("How many people does room 2.13 hold?")	"Is it warm?" = `missing` (implicit now, the fact is time-varying but no time is given): note it in F7
F3c criterion	a threshold, band or comparison standard is stated, or the question needs none	the question needs none ("What is the temperature?")	"too hot" = `missing`; "above 25 °C" = `specified`. **Unchanged in v2** — see the note above the table

F4 Evidence relation (single label)

Could the building's data, **if held**, establish the answer itself?

Value	Definition	Constructed example
`direct`	A datum is the asked property	"What was the CO_2 in room 5.01 at 10:00?"
`proxy`	Only a related quantity can exist; it cannot establish the asked property	"Is the room free?" when only occupancy is meant → presence ≠ booking. "Will I feel tired in there?" → CO_2 is a proxy for a feeling
`not_establishable`	No building datum can answer, even in principle	"Who is the best lecturer on floor 3?"
`unclear`	Cannot decide from the text	—

Rule. Prefer `direct` when a typical building could hold the asked property itself (bookings, capacity), even if this building does not. F4 is about the question, not about Building A.

F5 Answerable in principle (single label)

For a typically instrumented building with records (not Building A specifically):

`yes` (fully) · `partly` (some of the request) · `no`.

F6 Purpose stated (single label)

`yes` if the wording states **why** the asker wants it ("…so I can plan my meeting"); otherwise `no`. Code only what is stated.

F7 Coding certainty

`clear` · `uncertain` · `discuss`, plus a free-text note for anything not `clear`. Rows marked `discuss` by either coder go to the adjudication meeting first.

Procedure (plan 03 §3.4)

1. **Development session** (~2 h): the first author codes the 100 development rows, refines the definitions, briefs the second coder (30 minutes), and freezes v1 (SHA-256).
2. **Validation:** the first author codes all 400 validation rows; the second coder codes a seeded 100 of them (25 per stage, `../F_Analysis/draw_double_coding_sample.py`). Both work independently, blind to the automated labels and to each other. Use `sheets/partA_coding_sheet_TEMPLATE.csv`.
3. **External validation** (after the sessions): code the pooled participants' own questions mixed with 90 corpus rows, blind to source; the second coder codes a seeded 60 (`sheets/fresh_vs_corpus_coding_sheet_TEMPLATE.csv`).
4. **Agreement and adjudication** as in `Agreement_and_Adjudication_Procedure.md`.

Dates for each step are in the first author's own plan (the AI-made dates that stood here are withdrawn).

Changes made at the development session (v1)

Field	Change	Reason	Rows that prompted it
—	none: no development session was held (24 October 2026)	the submission date	—

Changes made in v2

Field	Change	Reason	Rows that prompted it
F3a place	added an explicit `n/a` rule (v1 left it undefined)	v1 defined `n/a` only for F3c; the second coder, coding blind, used no `n/a` at all for F3a across all 160 sampled items, which lowers measured agreement without reflecting a real disagreement	found checking the second coder's F3 codes against the first author's, 28 October 2026
F3b time	added an explicit `n/a` rule (v1 left it undefined)	as above	as above
F3c criterion	none — definition unchanged	the gap between coders here (9% vs 72% `specified` on the Part A sample) is not explained by a missing category; put to the second coder directly instead of patched unilaterally	as above